

CICD shell

# Table of Contents

Setup .....	1
Usage .....	1
>_ facts .....	3
>_ data .....	4
>_ runpuppet .....	4
>_ result .....	5
>_ service .....	5
>_ du .....	6
>_ setfacts .....	6
>_ console .....	6
>_ orch .....	7
>_ help .....	7
Authentication .....	7
Install outside the devbox .....	8
Changelog .....	8
v1.4, /05/2017 .....	8
v1.3, 08/05/2017 .....	8
v1.2, 26/04/2017 .....	9
v1.1, 19/04/2017 .....	9



# Setup

To use the `cicd shell`, you need to fill in the [configuration file](#) located in `/vagrant/config/shell`.

The shell is installed with the devbox. Please read the section *Install\_outside\_the\_devbox* if you need to install it on a different system.

# Usage

The name of the command line utility is `cicd`. The first mandatory position argument is the zone (dev, testing, staging or prod).

## *General scheme of the command line*

```
cicd ZONE facts [ROLE] [-n NODE] [-g GROUP] [-s STACK] [--all] [--down]
cicd ZONE data [-k KEY] [ROLE] [-n NODE] [-g GROUP] [-s STACK] [--all]
cicd ZONE runpuppet [ROLE] [-n NODE] [-g GROUP] [-s STACK]
cicd ZONE service ACTION SERVICE [ROLE] [-n NODE] [-g GROUP] [-s STACK]
cicd ZONE orch CMD [ROLE] [-n NODE] [-g GROUP] [-s STACK]
cicd ZONE result [-j JID] [-n NUM] [--raw]
cicd ZONE console

cicd doc (html | modules | mod MOD)
```

The auto-completion feature helps you to complete the command. You can also request the help at each level:

## General help output for cicd

→ cicd

CICD - command line utility (v1.4.0)

Usage: cicd (doc | ZONE (console | stats | data | orch | facts | ping | du | service | runpuppet | sync | setfacts | result | gentags))

Available options:

-h,--help	Show this help text
ZONE	ZONE (dev testing staging prod)

Available commands:

doc	Documentation utilities
console	Open the cicd console
stats	Stats (special permission required)
data	Return configuration data for a specific property
orch	Run an orchestration command on the infrastructure
facts	Return essential facts about nodes
ping	Ping nodes
du	Return disk usage
service	Service management for a specific node
runpuppet	Apply puppet configuration
sync	Syncmetavar data from master to nodes
setfacts	Set/update the 4 base machine facts
result	Display the results of the most recent jobs executed by the user or for a specific id
gentags	Generate node completion file

## Help output for facts

→ cicd staging facts -h

Usage: cicd facts [--down] [--all] [ROLE] [-n NODE] [-g GROUP] [-s STACK] [--raw] [-v|--verbose]

Return essential facts about nodes

Available options:

ROLE	Role name maybe prefixed by a subgroup ('subgroup.role')
-n NODE	Target node
-g GROUP	Target subgroup
-s STACK	Target stack/hostgroup
--down	Query down node
--all	Target whole the known stacks
--raw	Raw output (no jq)
-v,--verbose	Display the executed command
-h,--help	Show this help text



- Commands are executed remotely through an API. Behind the scene they call either the **puppetdb**, the **saltmaster** or the **pgserver**.
- Commands to the saltmaster together with their results are recorded in a centralized database included the date and name of the person that executes them.
- By default, all commands target a specific default hostgroup/stack defined in **/vagrant/conf/shell**

## >\_ facts

The command displays a subset of important facts (static information) about your nodes such as the **fqdn**, **ip**, **os**, **role**, ...

You can toggle the **facts** query to target all hostgroups/stacks with the **--all** flag. Here is how to get all facts for all slaves in every stack:

```
λ ~ → cied prod facts jenkins.slave --all
{
  "fqdn": "SVAPPCAVL595.cirb.lan",
  "ip": "192.168.34.153",
  "os": "CentOS 6.6",
  "hostgroup": "irisbox",
  "subgroup": "jenkins",
  "role": "slave"
  "puppet run": "Tue Apr 18 14:26:15 CEST 2017",
  "jenkins job": "633"
}
{
  "fqdn": "SVAPPCAVL649.prd.srv.cirb.lan",
  "ip": "192.168.34.9",
  "os": "RedHat 6.7",
  "hostgroup": "iam",
  "subgroup": "jenkins",
  "role": "slave"
  "puppet run": "Tue Apr 18 14:26:15 CEST 2017",
  "jenkins job": "633"
}
...
```

As usual, use **-n** to target a single node:

```
→ cicc prod facts -n svappcavl771.prd.srv.cirb.lan
{
  "fqdn": "svappcavl771.prd.srv.cirb.lan",
  "ip": "192.168.34.81",
  "os": "RedHat 7.2",
  "hostgroup": "fmx",
  "subgroup": "jenkins",
  "role": "slave",
  "puppet run": "Thu Jan 26 11:06:00 CET 2017",
  "jenkins job": "30"
}
```



Use the `--down` flag to gather **facts** on a disconnected minion.

## >\_ data

The command displays configuration data about your node. For instance you might display the docker version of your jenkins slave:

```
→ cicc prod data jenkins.slave -k docker::version
{
  "fqdn": "svappcavl736.cirb.lan",
  "subgroup": "jenkins",
  "role": "slave",
  "docker::version": "1.9.1-25.e17"
}
```

To display ALL known configurations for a specific node:

```
→ cicc prod data -n svappcavl771.prd.srv.cirb.lan
```

## >\_ runpuppet

The command runs the puppet agent on one or multiple nodes. When a node is specified with `-n`, the command will wait back for a result.

```
→ cicc dev runpuppet -n svappcavl000.dev.srv.cirb.lan
```

On all other cases, the command first asks for confirmation, then returns quickly with a **jobid**. The process is asynchronous because it might take quite a while to complete.

Here are some examples:

```
→ cicc dev runpuppet ①  
→ cicc dev runpuppet -g jenkins ②  
→ cicc dev runpuppet jenkins.slave ③
```

① run puppet on all the dev nodes of your stack

② run on a subgroup of machines

③ target a role

In a second step, you use `>_ result` to retrieve from the database the result of your call [1: polling is currently the sole supported workflow, server push notification could be implemented in the future].

## >\_ result

You can view the result of a `runpuppet` by using the provided job id (`jid`)

```
→ cicc testing result -j 20160621104434055991
```

In case the result is not yet available the command will automatically be retry 12 times (3 min).



The pretty printer is tailored to work on jobid coming from `>_ runpuppet`. For all other JIDs, you should add the `--raw` flag.

You can also ask for the last n executed commands:

```
→ cicc testing result -n 2
```

## >\_ service

To know if a service is up and running, you would use:

```
→ cicc prod service status docker jenkins.slave  
{  
  "svappcavl736.prd.srv.cirb.lan": true  
}
```

You can also restart a service. However such operation is only allowed for a single machine. Here is how to restart the `nexus` service :

```
→ cird prod service restart nexus -n svappcavl761.prd.srv.cirb.lan
{
  "svappcavl761.prd.srv.cirb.lan": true
}
```

## >\_ du

The command displays disk usage. Try:

```
→ cird staging du -n svappcavl703.sta.srv.cirb.lan
```

## >\_ setfacts

To set (or update) the four basic facts on a specific machine:

```
→ cird dev setfacts -n fqdn --subgroup jenkins --role slave --zone dev --hostgroup
bas
```

You can of course update just one fact with:

```
→ cird dev setfacts -n fqdn --subgroup jenkins2
```



the **setfacts** subcommand always requires a target node (-n)

## >\_ console

For longer session within a specific zone, you can save some typing by opening a **console** for that zone. Inside the console, you would omit the zone from the command line. Here is an example:

```
→ cird staging console

[cird prod]$ facts
```

Another usage of the console is to run specific **salt** commands that are not exposed by the **cird** command line. This is done via the pep shortcut. For instance:

```
$ pep -G 'hostgroup:iam' file.replace '/etc/resolv.conf' pattern='192.168.34.250' repl='192.168.34.244' ①

$ pep -L fqdn1,fqdn2 --client=local_async puppetutils.run_agent ②
```

① -G means **grain** target (*grains* is the salt terminology for facts).



② -L means **list** target

local\_asyn means the command is asynchronous and does not display its result (just a jid)



- Have a look at the saltstack documentation to learn more about [targeting minions](#).
- Take a look [here](#) for a list of possible commands.

## >\_ **orch**

Salt is able to orchestrate deployment scenarios across machines.

The orchestration is executed on the salt master to allow inter minion requisites, like ordering the application of states on different minions that must not happen simultaneously, or for halting the state run on all minions if a minion fails one of its states (more about this topic can be found [in the saltstack website](#)).

To write some specific orchestration scripts for your stack, you need to request a salt stack repository. For **bos** it would be named **salt-stack-bos**. This process is similar to the creation of **puppet-stack-bos**. The scripts should sit in the **orch** folder. You can find some examples [here](#).

Orchestrate commands are executed with:

```
→ cird testing orch CMD
```

## >\_ **help**

The **help** subcommand will open the guide in a browser, display the list of available salt module and show the help for each of them.

```
→ cird doc
Usage: cird doc (html | modules | mod)
Documentation utilities

Available options:
  -h, --help          Show this help text

Available commands:
  html                Open the documentation in a browser
  modules             Output all possible salt execution modules
  mod                 Doc about a specific salt module
```

# Authentication

The permissions to target machines and perform actions are realized through our Active directory. As an example to access the machines of the `middleware` hostgroup, you will need to be part of the `GP_APP_SALT_MIDDLEWARE` group.

These permissions should have been set for you already. If they don't, please contact the `cicd` team.

## Install outside the devbox

Before installing the `cicd-shell` on any linux system [2: `macos` might also work], you will need:

1. the `nix package manager` installed and active for your user.
2. the `cirb nixpkgs config`

You can then proceed to install with:

```
nix-env -f ~/.config/nixpkgs/pin.nix -i cicd-shell ①
```

① the `-f` flag ensures that we point to the same nixpkgs version but can be omitted

If you haven't installed `nix` already, here is the quick how to:

```
bash <(curl https://nixos.org/nix/install)
```

This will perform a single-user installation of Nix, meaning that `/nix` is owned by the invoking user. The script will only invoke `sudo` to create `/nix` if it doesn't already exist. At that point, the script will prompt you for a password.

To activate `nix` in your shell, add the following line in your `.bash_profile`:

```
source ~/.nix-profile/etc/profile.d/nix.sh'
```

## Changelog

### v1.4, /05/2017

- add `cicd ZONE setfacts --subgroup= --role=server --hostgroup=middleware

### v1.3, 08/05/2017

## changes

- `cicd help` is now `cicd doc`

## v1.2, 26/04/2017

- add
  - `cicd help html`
  - `cicd help mod MOD`
  - `cicd help modules`
- add `--raw` option to display output without `jq` pretty printer
- add `--verbose` option to display the executed command
- allow `--all` to be used with the `data` subcommand
- increase default timeout (up to 3min) for synchronous `runpuppet`
- provide completion for salt modules in the console

## v1.1, 19/04/2017

- `retries` when a command fails
- fix issues with the `result` subcommand output